

## Question 1: Data understanding

- **Read data:**

Import modules and read the three csv files with pandas. The three data frames `user_features.csv`, `product_features.csv`, and `click_history.csv` are named as `userFeatures`, `proFeatures` and `click`, respectively

- **Check the structure of data**

**userFeatures:** has 12000 rows and 4 columns, with `user_id` as integer, `number_of_clicks_before` as object, `ordered_before` as boolean, `personal_interest` as object, and it has many values in a cell. It needs to be separated in the further steps.

```
# read user features data and check the data structure
userFeatures=pd.read_csv("user_features.csv")
userFeatures.info() # check the structure of the data
userFeatures.head() # check what the data looks like
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12000 entries, 0 to 11999
Data columns (total 4 columns):
user_id                12000 non-null int64
number_of_clicks_before 11500 non-null object
ordered_before         12000 non-null bool
personal_interests     12000 non-null object
dtypes: bool(1), int64(1), object(2)
memory usage: 293.1+ KB
```

	user_id	number_of_clicks_before	ordered_before	personal_interests
0	104939	2	True	['body', 'makeup', 'nail', 'hand', 'foot', 'me...
1	101562	2	True	['men_skincare', 'men_fragrance', 'tools', 'sk...
2	102343	2	True	['tools', 'makeup', 'foot', 'nail']
3	106728	5	True	['hand', 'men_skincare']
4	107179	0	True	['makeup', 'body', 'skincare', 'foot', 'men_sk...

**proFeatures:** has 1000 rows and 5 columns, with `product_id` as integer, category as object, `on_sale` as boolean, `number_of_reviews` as integer, `avg_review_score` as float.

```
# read product features data and check the data structure
proFeatures=pd.read_csv("product_features.csv")
proFeatures.info() # check the data structure
proFeatures.head() # check the head of the dataset
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 5 columns):
product_id            1000 non-null int64
category              1000 non-null object
on_sale               1000 non-null bool
number_of_reviews     1000 non-null int64
avg_review_score      1000 non-null float64
dtypes: bool(1), float64(1), int64(2), object(1)
memory usage: 32.4+ KB
```

	product_id	category	on_sale	number_of_reviews	avg_review_score
0	1134	tools	False	101	3.349452
1	1846	skincare	False	111	5.000000
2	1762	fragrance	False	220	4.882706
3	1254	hair	True	446	5.000000

**click:** has 35990 rows and 3 columns, with `user_id` as integer, `product_id` as integer, `clicked` as boolean.

```
# read product features data and check the data structure
click=pd.read_csv("click_history.csv")
click.info() # check the structure of the dataframe
click.head() # take a look at the head of the data
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 35990 entries, 0 to 35989
Data columns (total 3 columns):
user_id      35990 non-null int64
product_id   35990 non-null int64
clicked      35990 non-null bool
dtypes: bool(1), int64(2)
memory usage: 597.6 KB
```

	user_id	product_id	clicked
0	104863	1350	False
1	108656	1321	True
2	100120	1110	False
3	104838	1443	True
4	107304	1397	True

- **Check the descriptive statistics of data**

**userFeatures:** the count of user\_id=12000, informing there are 12000 users, other descriptive statistics is not informative since this variable is only the user's identification.

```
#get the summary of the user features
userFeatures.describe() #
```

	user_id
count	12000.00000
mean	105999.50000
std	3464.24595
min	100000.00000
25%	102999.75000
50%	105999.50000
75%	108999.25000
max	111999.00000

**proFeatures:** For product\_id: count=1000, indicating 1000 products. For number\_of\_reviews, minimum=66 and maximum=2307390, and it can be seen there is a big difference. For avg\_review\_score, it has a minimum of -1 and maximum of 5.

```
#Check the summary of the product features
proFeatures.describe() # Check the descriptive sta
```

	product_id	number_of_reviews	avg_review_score
count	1000.000000	1.000000e+03	1000.000000
mean	1499.500000	1.157725e+05	2.660656
std	288.819436	5.028997e+05	1.741875
min	1000.000000	6.600000e+01	-1.000000
25%	1249.750000	2.570000e+02	1.428969
50%	1499.500000	4.710000e+02	2.769397
75%	1749.250000	7.042500e+02	4.180860
max	1999.000000	2.307390e+06	5.000000

**Click:** The count of user\_id and product\_id is 35990. The other descriptive statistics are not informative since they are only id numbers.

```
#check the summary of the click history
click.describe() # Check the description
```

	user_id	product_id
count	35990.000000	35990.000000
mean	106017.080161	1500.232898
std	3483.480090	288.101984
min	100001.000000	1000.000000
25%	102976.500000	1250.000000
50%	106060.000000	1503.000000
75%	109049.000000	1749.000000
max	111999.000000	1999.000000

## Question 2: Data cleaning and preprocessing

- **Check duplicated rows:**

The output shows there is no duplicated rows in each of the three data frames.

```
#check duplicated rows
print(userFeatures.duplicated().sum()) #
print(proFeatures.duplicated().sum()) #
print(click.duplicated().sum()) # output

0
0
0
```

- **Check missing value:**

**uerFeatures:** 500 missing values in the column number\_of\_clicks\_before

**proFeatures:** no missing values

**click:** no missing values

```
# check missing value

print("Missing value for userFeatures: \n", userFeatures.isnull().sum())
# the variable number_of_clicks_before has 500 missing values

print("\n Missing value for proFeatures: \n", proFeatures.isnull().sum().values.sum())

print("\n Missing value for clicks: \n", click.isnull().sum().values.sum()) # no missing values

Missing value for userFeatures:
  user_id      0
number_of_clicks_before    500
ordered_before      0
personal_interests      0
dtype: int64

Missing value for proFeatures:
0

Missing value for clicks:
0
```

- **Handling missing values**

**userFeatures:** I checked the rows with missing values, it looks the missing values are randomly distributed and it is hard to tell what reason caused the missing values. And then I checked the distribution of `number_of_clicks_before`, and found the `number_of_clicks_before=2` is the most frequent number. Therefore, I decided to fill the missing value with 2.

```
#check the distribution of the number_of_clicks_before
print(userFeatures.groupby(["number_of_clicks_before"]).count()) # the
#so I decide to replace the missing value with 2
userFeatures.fillna(2,inplace=True) # fill the missing value with 2
```

	user_id	ordered_before	personal_interests
number_of_clicks_before			
0	2196	2196	2196
1	2399	2399	2399
2	2710	2710	2710
3	59	59	59
4	711	711	711
5	1057	1057	1057
6+	2368	2368	2368

```
#check if there is still missing value
userFeatures.isnull().sum() # ouput is 0, no missing value
```

```
user_id      0
number_of_clicks_before  0
ordered_before  0
personal_interests  0
dtype: int64
```

- **Separate cells where there are multiple values, get\_dummies to convert them to numeric variables and check missing value:**

**userFeatures:** I found the column `personal_interest` has many values in a cell and this information seems important since it includes the interest of the users and it might be helpful in predicting if some products will be clicked or not. Therefore, I decided to separate the values and get dummies to convert them into numeric variables.

Drop the original `personal_interest` column, since the dummies has replaced its information.

Again, check the missing values, and I found there are 1096 missing values in those interest variables. I think these missing values are from those cells where the user does not have that interest, so I fill these missing values with 0.

```
# separate personal_interests variable
import ast
```

```
userFeatures.personal_interests = userFeatures.personal_interests.apply(lambda x: ast.literal_eval(x))
```

```
# convert it to numerical variable using get_dummies
```

```
dumPerInt = pd.get_dummies(userFeatures.personal_interests.apply(pd.Series).stack()).sum(level=0)
```

```
dumPerInt.head(3)
```

	body	foot	fragrance	hair	hand	makeup	men_fragrance	men_skincare	nail	skincare	tools
0	1	1	1	1	1	1	1	0	1	0	1
1	1	0	0	0	0	1	1	1	1	1	1
2	0	1	0	0	0	1	0	0	1	0	1

```
# Concatenate the original dataframe with the dummy dataframe
```

```
userFeatures1 = pd.concat([userFeatures, dumPerInt], axis=1)
```

```
userFeatures1.head(3)
```

	user_id	number_of_clicks_before	ordered_before	personal_interests	body	foot	fragrance	hair	hand	makeup	men_fragrance	men_skincare	nail	skinc
0	104939		2	True	[body, makeup, nail, hand, foot, men_fragrance...]	1.0	1.0	1.0	1.0	1.0	1.0	0.0	1.0	
1	101562		2	True	[men_skincare, men_fragrance, ...]	1.0	0.0	0.0	0.0	1.0	1.0	1.0	1.0	

```
# drop the personal_interests column since the dummy vairables has
userFeatures1.drop(["personal_interests"], axis=1, inplace=True)
```

```
# check if there is new missing values appeared
userFeatures1.isnull().sum() # all the dummies has 1096 missing \
```

```
user_id      0
number_of_clicks_before  0
ordered_before  0
body      1096
foot      1096
fragrance  1096
hair      1096
hand      1096
makeup     1096
men_fragrance  1096
men_skincare  1096
nail       1096
skincare    1096
tools      1096
dtype: int64
```

```
# fill the NA with 0 since the NAs ar
```

```
userFeatures1.fillna(0, inplace=True)
```

```
# again check the missing values
userFeatures1.isnull().sum()
```

```
user_id      0
number_of_clicks_before  0
ordered_before  0
body      0
foot      0
fragrance   0
hair      0
hand      0
makeup     0
men_fragrance  0
men_skincare  0
nail       0
skincare    0
tools      0
dtype: int64
```

- **Merge data frames:**

Merge **userFeatures** with **click** on “user\_id” and named as **df1**, and then merge the new dataframe **df1** with **proFeatures** on “product\_id” named as **df2**. Check the missing values and shape. Now **df2** has no missing values and has 35990 rows and 20 columns.

```
# merge the data frame
df1=pd.merge(userFeatures1,click,on="user_id")
df1.head()
df2=pd.merge(df1,proFeatures,on="product_id")
df2.isnull().sum() # there is no missing value
df2.shape
```

```
(35990, 20)
```

- **Encoding the category variable into numeric variables:**

Encoding columns “ordered\_before, category, clicked, on\_sale, number\_of\_clicks\_before” into numeric variables.



```
# encoding the category variables into numbers
df2.category.unique() # check the category variable

df2.category=df2.category.astype("category").cat.codes #encoding the category variable

# Encoding the other categorical variables
df2.ordered_before=df2.ordered_before.astype("category").cat.codes #encoding ordered_before

df2.clicked=df2.clicked.astype("category").cat.codes #encoding_clicked

df2.on_sale=df2.on_sale.astype("category").cat.codes # encoding on_sale

df2.number_of_clicks_before=df2.number_of_clicks_before.astype("category").cat.codes #encoding number_of_clicks before

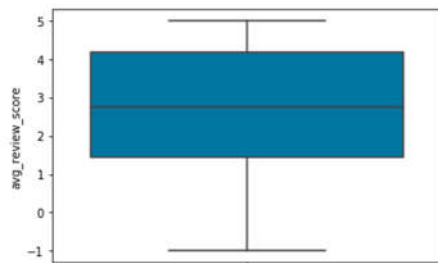
df2.head()
```

	user_id	number_of_clicks_before	ordered_before	body	foot	fragrance	hair	hand	makeup	men_fragrance	men_skincare	nail	skincare	tools	pr
0	104939	3	1	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.0	1.0	0.0	1.0	
1	101992	2	1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	

- **Check outliers:**

Based on the previous understanding of the dataframes. The variables **avg\_review\_score** and **number\_of reviews** are informative numerical variables.

```
# make plot to check the outliers for avg_review_score
sns.boxplot(y=df2.avg_review_score) # from the boxplot, i
<matplotlib.axes._subplots.AxesSubplot at 0x153bf1b4b70>
```



Therefore, I only checked the outlier of those two variables. For **avg\_review\_score**, from the boxplot, there is no outliers if the outliers are defined as more than 1.5IQR and less than -1.5IQR.

Also, I write a function to check the outliers based on  $\pm 1.5\text{IQR}$ , and found there is indeed no outliers in **avg\_review\_score**.

```
# define a function to check outliers
def outliers_IQR(Var):
    Q1, Q3 = np.percentile(Var, [25, 75]) # get Q1 and Q3 of a
    IQR = Q3-Q1 # calculate the IQR
    lower_bound=Q1-(IQR*1.5) # define the lower threshold of a
    upper_bound=Q3 + (IQR*1.5) # define the upper threshold of
    return np.where((Var > upper_bound) | (Var < lower_bound))

print(outliers_IQR(df2.avg_review_score)) # call the function a

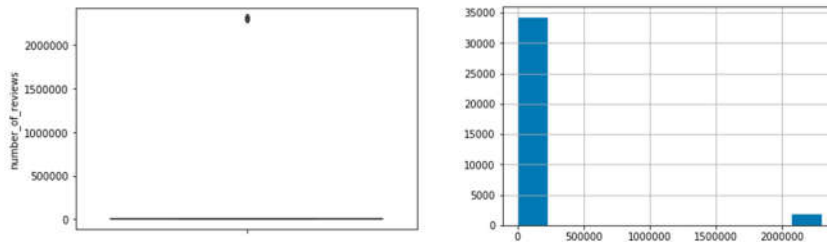
# output is empty array, indicating no outliers in this variable

(array([], dtype=int64),)
```

For **number\_of\_reviews**, from the boxplot and bar plot, there is a group of extremely big values. Also, I used the function to check the outliers and found there are 1770 outliers.

```
# visualize the distribution of number_of_reviews
sns.boxplot(y=df2.number_of_reviews) # it seems there is
<matplotlib.axes._subplots.AxesSubplot at 0x153bf14e160>

# check the histogram
df2.number_of_reviews.hist() # the number_of_reviews dist
<matplotlib.axes._subplots.AxesSubplot at 0x153bf0d2160>
```



```
# check the outliers using the function
review_Out=outliers_IQR(df2.number_of_reviews)
print(review_Out) # display the index of outliers
print(len(review_Out[0])) # display number of outliers 1770 outliers

(array([ 413,  414,  415, ..., 35859, 35860, 35861], dtype=int64),)
1770
```

- **Handling outliers**

I checked the distribution of outliers grouped by the “clicked” variable and found most (1126) of the outliers belong to the group clicked=0. According to a discussion with my previous applied statistics professor, if the outliers belong to a certain group, their existing is reasonable and it will not influence the result. Here most of the outliers belong to a group, I think it will not influence the result much, but I am still curious how much it will influence the result.

#### 2.6.1 keep the outliers---df2

```
df3=df2.loc[review_Out[0],:] # extracting the rows with outliers from the dataframe

df3.groupby("clicked").count() # it seems the most of outliers are from product have been clicked, so it mostly
# belongs to one group, I think maybe it will not influence the result much, I decide to keep the outliers
```

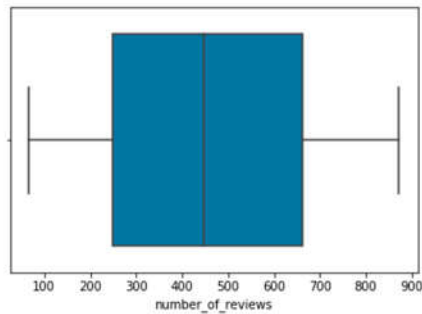
	user_id	number_of_clicks_before	ordered_before	body	foot	fragrance	hair	hand	makeup	men_fragrance	men_skincare
clicked											
0	1126	1126	1126	1126	1126	1126	1126	1126	1126	1126	1126
1	644	644	644	644	644	644	644	644	644	644	644

Therefore, I decide to **keep the outliers in df2**, and **make a copy of df2, which is df2a**. In **df2a**, I **removed the outliers** in order to compare the result from the two data frames in the further steps.

### 2.6.2 Remove the outliers---df2a

```
# removing outliers and save it to another dataframe to compare
df2a=df2.copy() # copy the merged dataset
df2a.drop(df2a.index[review_Out[0]],axis=0,inplace=True) #
sns.boxplot(df2a.number_of_reviews) #no more outliers anymore
print(df2a.shape)
```

(34220, 20)



## Question 3: Model building and evaluating

- **Import modules and split the data**

Drop the variables “user\_id, product\_id, click” from df2 and df2a respectively, and take remaining data as features. Select the column “clicked” as target.

Split the features and target into training and testing data at the ratio of 0.7 to 0.3 from df2 and df2a respectively.

```
from sklearn.model_selection import train_test_split # import split module

#split the dataframe into features and target
X=df2.drop(columns=["user_id","product_id","clicked"],axis=1) # drop the three columns and set
Y=df2.loc[:, "clicked"] # set clicked column as target

#do the something as above using dataset without outliers
X1=df2a.drop(columns=["user_id","product_id","clicked"],axis=1)
Y1=df2a.loc[:, "clicked"]

# split the data with 0.7/0.3 for training and testing
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, random_state=0, train_size=0.7)

X1_train, X1_test, Y1_train, Y1_test = train_test_split(X1, Y1, random_state=0, train_size=0.7)
```

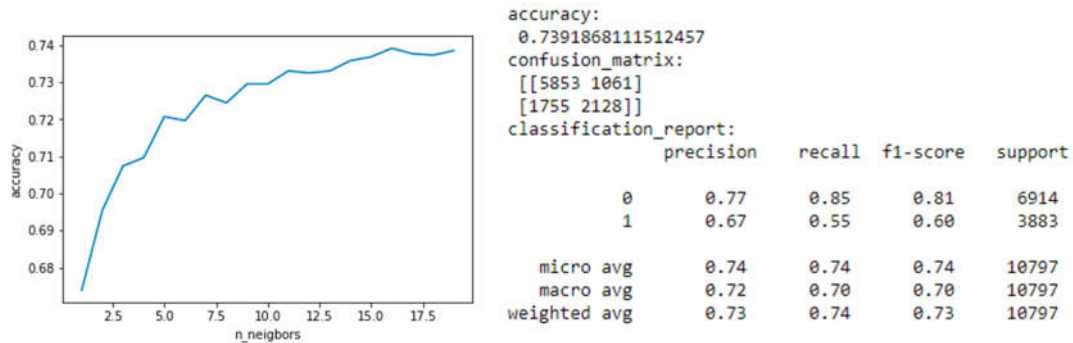
- **KNN (KNeighborsClassifier)**

Fit the KNeighborsClassifier model with training data and use the model to predict the test data.

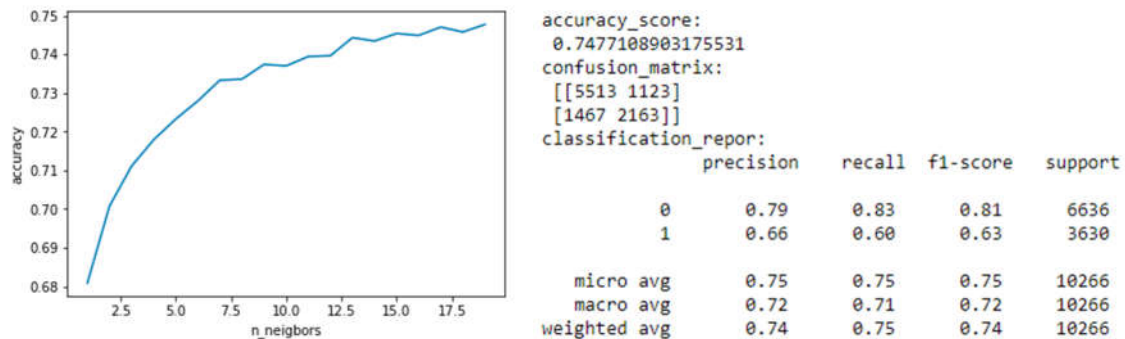
Check the accuracy with different number of n\_neighbors, k and choose the model with best accuracy to check the result.



For df2 (keep the outliers), it has the highest accuracy of 0.7391 when k=16. The precision and recall for the class 0 are 0.77 and 0.85 respectively. Precision and recall are also important for making a business decision since accuracy is sometimes miss-leading and decision should be made based on the consequence of the predicted result.



For df2a (removing the outliers), it has the highest accuracy of 0.7477 when k=19. The precision of class 0 is 0.79 and 0.83 respectively, the precision is a little bit better than data from df2, but recall is lower.

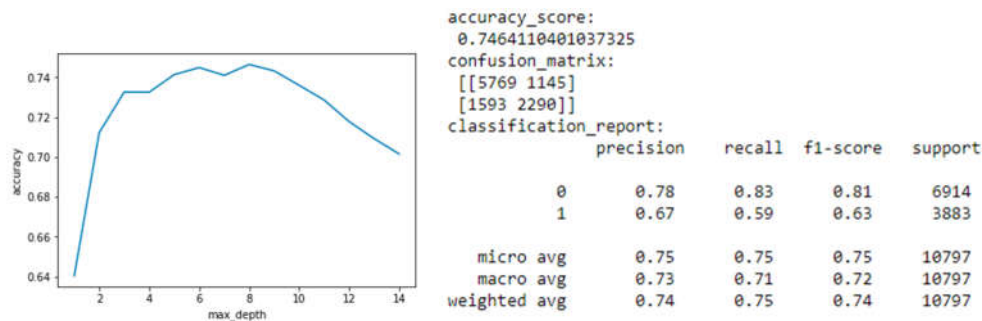


- **Decision Tree**

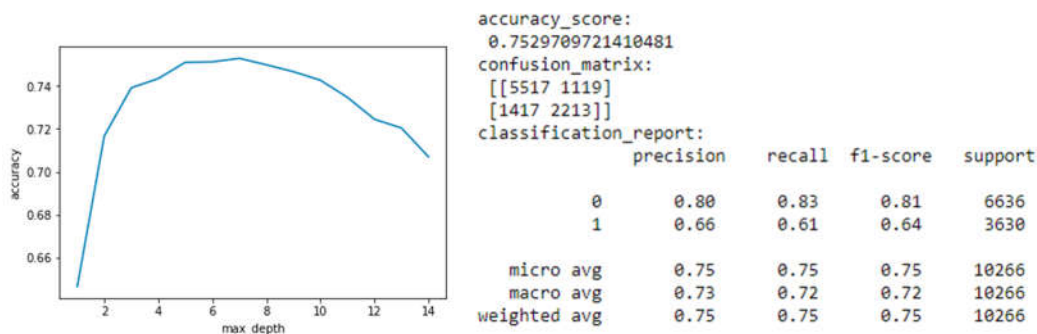
Fit the decision tree model with training data and use the model to predict the test data.

Check the accuracy with different number of max\_depth, k2 and choose the model with best accuracy to check the result.

For df2 (keep the outliers), it has the highest accuracy of 0.7464 when k2=8. The accuracy is better than KNeighborsClassifier model. But the precision is similar and recall is lower than KNeighborsClassifier model for df2.



For df2a (removing the outliers), it has the highest accuracy of 0.7530 when k2=7. The accuracy is better than data from df2. Also, the accuracy is higher than df2a fitted with KNeighborsClassifier model. Precision is the highest compared to the previous models. and recall is close to other models.



- **Random Forest**

Fit the random forest model (n\_estimator=30) with training data and use the model to predict the test data.

For df2 (keep the outliers), it has the highest accuracy of 0.7316. This model has lower accuracy than model fitted by KNeighborsClassifier and decision tree. The precision is also lower, but it has the high recall, which is also important for making business decisions.

```

accuracy_score:
0.7315921089191442
confusion_matrix:
[[5842 1072]
 [1826 2057]]
classification_report:

```

	precision	recall	f1-score	support
0	0.76	0.84	0.80	6914
1	0.66	0.53	0.59	3883
micro avg	0.73	0.73	0.73	10797
macro avg	0.71	0.69	0.69	10797
weighted avg	0.72	0.73	0.72	10797

For df2a (removing the outliers), it has the highest accuracy of 0.7358. The performance is better than df2 both in accuracy and precision and recall. It has lower accuracy than the other two models but high recall.

```
accuracy_score:
0.7358270017533606
confusion_matrix:
[[5553 1083]
 [1629 2001]]
classification_report:
      precision    recall  f1-score   support

     0       0.77      0.84      0.80      6636
     1       0.65      0.55      0.60      3630

 micro avg       0.74      0.74      0.74     10266
 macro avg       0.71      0.69      0.70     10266
 weighted avg     0.73      0.74      0.73     10266
```

**Question 4: Which model has the best performance? What have you learned from the models you built?**

#### 1. Which model has the best performance?

- The accuracy of each model is in the table below:

	df2 (keep outliers)	df2a (removing outliers)
K Neighbors Classifier	0.7391	0.7477
Decision tree	<b>0.7464</b>	<b>0.7530</b>
Random Forest	0.7316	0.7358

- The decision tree model has the best performance for both datasets.
- Removing the outliers improves the model and get better accuracy than data with outliers.

#### 2.What we have learned from the model:

- Data cleaning and processing is time consuming, but it is very important and worth doing, since it is obvious that outliers influence the result of the prediction.
- Keep curious about models built from data with different preprocessing help gaining insight about how to improve model and diagnose problems in the future.
- Random forest model is based on decision tree and it should have got better performance, but more time is needed to improve the model the get a better performance and this work will be done in my future study.
- Although decision tree has the highest accuracy and precision, the recall is not the highest in the model. When we choose a model in real life, we should also consider the recall since it is relevant to the consequence of the problem.